

Visualizing high dimensional TEDS data

```
In [2]: import matplotlib.pyplot as plt
import pandas as pd

import seaborn as sns

import numpy as np
import scipy.stats as ss

from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
```

```
In [7]: df=pd.read_csv('__TEDSCountyLevel2008-2017_2.csv')
```

```
In [3]: %matplotlib inline
```

```
In [8]: df.head(10)
```

Out[8]:

	Alcohol	Amphetamine	Barbiturates	Benzine	Age	State	Cocaine	County Name	Gender	Hal	...	Pr
0	0	0	0	0	18-24	Indiana	0	Pulaski	Male	0	...	0
1	1	0	0	0	55+	Indiana	0	Pulaski	Male	0	...	0
2	1	0	0	0	18-24	Indiana	0	Pulaski	Male	0	...	0
3	0	0	0	0	55+	Indiana	0	Pulaski	Male	0	...	0
4	0	0	0	0	55+	Indiana	0	Pulaski	Male	0	...	0
5	1	0	0	0	25-34	Indiana	0	Pulaski	Male	0	...	0
6	1	0	0	0	35-44	Indiana	0	Pulaski	Male	0	...	0
7	1	0	0	0	25-34	Indiana	0	Pulaski	Male	0	...	0
8	1	0	0	0	0-18	Indiana	0	Pulaski	Male	0	...	0
9	1	0	0	0	25-34	Indiana	0	Pulaski	Male	0	...	0

10 rows × 50 columns

```
In [9]: #Limit the number of features
features = ['Year', 'County Name', 'Gender', 'Race', 'Age', 'Alcohol', 'Heroin', 'Marijuana', 'Opioid', 'Cocaine', 'Other Drugs', 'Pain Killers', 'Methamphetamine', 'Synthetic Drugs']
df_subset = df[features]
df_subset.head()
```

Out[9]:

	Year	County Name	Gender	Race	Age	Alcohol	Heroin	Marijuana	Opioid	Cocaine	Other Drugs	Pain Killers	M
0	2009	Pulaski	Male	White	18-24	0	0	1	0	0	0	0	0
1	2009	Pulaski	Male	White	55+	1	0	0	0	0	0	0	0
2	2009	Pulaski	Male	White	18-24	1	0	1	0	0	0	0	0
3	2009	Pulaski	Male	White	55+	0	0	1	0	0	0	0	0
4	2009	Pulaski	Male	White	55+	0	0	1	0	0	0	0	0

```
In [10]: import altair as alt
alt.renderers.enable('notebook')
```

Out[10]: `RendererRegistry.enable('notebook')`

```
In [11]: #Create a dataset for the 9 substances to use in correlations
features = ['Alcohol', 'Heroin', 'Marijuana', 'Opioid', 'Cocaine', 'Other Drugs', 'Pain Killers', 'Methamphetamine', 'Synthetic Drugs']
df_only_features = df_subset[features]
df_only_features.head()
```

Out[11]:

	Alcohol	Heroin	Marijuana	Opioid	Cocaine	Other Drugs	Pain Killers	Methamphetamine	Synthetic Drugs
0	0	0	1	0	0	0	0	0	0
1	1	0	0	0	0	0	0	0	0
2	1	0	1	0	0	0	0	0	0
3	0	0	1	0	0	0	0	0	0
4	0	0	1	0	0	0	0	0	0


```
In [12]: n_samples, n_features = df_only_features.shape
print(n_samples)
print(n_features)
```

335866
9

```
In [11]: #Check correlation between substances
import seaborn as sns

f, ax = plt.subplots(figsize=(9, 9))
corr = df_only_features.corr()
sns.heatmap(corr, mask=np.zeros_like(corr, dtype=np.bool), cmap=sns.diverging_palette(220, 10, as_cmap=True), center=0,
            square=True, ax=ax, annot=True)
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x1bd9f9260b8>
```



There is strong correlation between Pain Killers and Synthetic Drugs, supported by the Opioids. Heroin appears in 57% of the cases with Opioids. Alcohol is slightly negatively correlates with all four of them.

Remarkably, there appears to be no appreciable correlation between opioid use and such commonly accepted and complementary socio-economic markers as alcohol, marijuana, or methamphetamine use. In fact, these three factors are likely to be biased towards older population (alcohol), younger population (marijuana) and underprivileged rural population (methamphetamine).

Thus, it appears that opioid abuse have roots that are different from the common socio-economic maladies reflected in these three markers.

```
In [13]: df_subset['County Name'].nunique()
```

```
Out[13]: 93
```

```
In [13]: df_subset['Year'].nunique()
```

```
Out[13]: 10
```

```
In [14]: # Total number of records  
len(df_subset)
```

```
Out[14]: 335866
```

```
In [15]: #How many records per year  
df_subset['Year'].value_counts()
```

```
Out[15]: 2017    37459  
        2015    34739  
        2011    34677  
        2014    34470  
        2012    33984  
        2013    33616  
        2016    33170  
        2009    32049  
        2010    31827  
        2008    29875  
        Name: Year, dtype: int64
```

```
In [15]: #How many records per county  
df_subset['County Name'].value_counts()
```

```
Out[15]: Marion      41678
         Lake        25329
         Allen       16940
         Vanderburgh 14085
         Monroe      13842
         Saint Joseph 13359
         Delaware    10453
         Madison     9594
         Elkhart     8309
         Hamilton    8033
         Vigo        7895
         Howard      6235
         Bartholomew 6037
         Porter      5455
         LaPorte     5086
         Morgan      5057
         Lawrence    4557
         Tippecanoe  4514
         Wayne       4391
         Grant       4241
         Dearborn    3981
         Clark       3857
         Hendricks   3488
         Knox        3350
         Noble       3024
         Montgomery  3006
         Warrick     2969
         Kosciusko   2946
         Jefferson   2888
         Johnson     2834
         ...
         Perry       1407
         Rush        1281
         Shelby      1246
         Vermillion  1241
         Wells       1195
         Jay         1193
         Orange      1174
         Franklin    1115
         Huntington  1107
         Parke       1084
         Whitley     1031
         Carroll     967
         Brown       922
         Pulaski     906
         Jasper      902
         Sullivan    812
         Blackford   759
         Washington  681
         Fountain    662
         Martin      638
         Tipton      563
         Switzerland 562
         Crawford    511
         Harrison    461
         Pike        433
         Benton      373
         Newton      315
         Ohio        293
         Union       284
```

Warren 229
Name: County Name, Length: 93, dtype: int64

```
In [16]: # General method to variably group the data
def group_dataset(df, col1, col2):
    df_grouped=df.groupby([col1,col2])
    return df_grouped
```

```
In [17]: #Group by Year and County
df_grouped=group_dataset(df_subset, 'Year','County Name')

#Excluded pain killers and synthetic drugs as highly correclated with opioids
feature=['Alcohol','Heroin','Marijuana','Opioid','Cocaine','Other Drugs','Methamphetamin
e']
df_grouped=df_grouped[feature].mean()

df_grouped.head()
```

Out[17]:

		Alcohol	Heroin	Marijuana	Opioid	Cocaine	Other Drugs	Methamphetamine
Year	County Name							
2008	Adams	0.905263	0.021053	0.684211	0.105263	0.231579	0.0	0.010526
	Allen	0.863495	0.012480	0.641186	0.048362	0.325273	0.0	0.017161
	Bartholomew	0.651106	0.004914	0.542998	0.226044	0.255528	0.0	0.299754
	Benton	0.840000	0.000000	0.520000	0.040000	0.160000	0.0	0.040000
	Blackford	0.833333	0.000000	0.620370	0.157407	0.175926	0.0	0.009259

```
In [18]: #Group by County
df_grouped_c=group_dataset(df_subset,'County Name', 'Year')

#exclude pain killers and synthetic drugs as highly correclated with opioids
feature=['Alcohol','Heroin','Marijuana','Opioid','Cocaine','Other Drugs','Methamphetamin
e']
df_grouped_c=df_grouped_c[feature].mean()

df_grouped_c.head()
```

Out[18]:

		Alcohol	Heroin	Marijuana	Opioid	Cocaine	Other Drugs	Methamphetamine
County Name	Year							
Adams	2008	0.905263	0.021053	0.684211	0.105263	0.231579	0.000000	0.010526
	2009	0.877863	0.007634	0.549618	0.099237	0.175573	0.000000	0.030534
	2010	0.613497	0.141104	0.496933	0.300613	0.184049	0.147239	0.042945
	2011	0.655405	0.141892	0.540541	0.331081	0.155405	0.108108	0.060811
	2012	0.756098	0.073171	0.528455	0.203252	0.178862	0.138211	0.040650

```
In [19]: # Group by Year to genralize the output
df_grouped_year1=df_subset.groupby('Year')
df_grouped_year1=df_grouped_year1[feature].mean()

df_grouped_year1.head()
```

Out[19]:

	Alcohol	Heroin	Marijuana	Opioid	Cocaine	Other Drugs	Methamphetamine
Year							
2008	0.661891	0.038226	0.506644	0.155213	0.223397	0.000000	0.087464
2009	0.682798	0.049237	0.514837	0.177821	0.194359	0.000000	0.091235
2010	0.632168	0.055833	0.490998	0.186006	0.157602	0.119553	0.096962
2011	0.606050	0.068662	0.483837	0.220607	0.146899	0.138593	0.103556
2012	0.590837	0.094191	0.471398	0.275012	0.156809	0.154926	0.112612

```
In [20]: # General method to pivot by two predictors
def create_pivot_df(df, col_name, index_name, val_name):
    pivot_df = df.pivot_table(columns=col_name, index= index_name, values= val_name, agg
func='mean')
    return pivot_df
```

```
In [21]: #create a pivot table for alcohol
create_pivot_df(df_grouped,col_name='Year',index_name='County Name', val_name='Alcohol')
.head()
```

Out[21]:

Year	2008	2009	2010	2011	2012	2013	2014	2015	2
County Name									
Adams	0.905263	0.877863	0.613497	0.655405	0.756098	0.812030	0.805195	0.697143	0.559
Allen	0.863495	0.834094	0.812057	0.794048	0.672963	0.736813	0.729648	0.674190	0.665
Bartholomew	0.651106	0.605505	0.442708	0.400815	0.458841	0.522551	0.501475	0.410745	0.345
Benton	0.840000	0.850000	0.684211	0.794118	0.852941	0.659091	0.708333	0.750000	0.780
Blackford	0.833333	0.787234	0.216216	0.523810	0.519231	0.386364	0.440678	0.454545	0.378

```
In [22]: def draw_heatmap(df, nrows, ncols):
    f, ax = plt.subplots(figsize=(nrows, ncols))
    sns.heatmap(df, cmap=sns.cubehelix_palette(8, start=.5, rot=-.75),ax=ax,xticklabels
=2, square=True, vmin=0, vmax=.9)
```

```
In [23]: def draw_heatmap_corr(df, nrows, ncols):
    f, ax = plt.subplots(figsize=(nrows, ncols))
    df=df.corr()
    sns.heatmap(df, cmap=sns.diverging_palette(220, 10, as_cmap=True), center=0,ax=ax,
square=True,annot=True,cbar_kws={"shrink": .5})
```




```
In [24]: #Transpose the matrix
df_grouped_year_T= df_grouped_year1.T
df_grouped_year_T
```

Out[24]:

Year	2008	2009	2010	2011	2012	2013	2014	2015
Alcohol	0.661891	0.682798	0.632168	0.606050	0.590837	0.585792	0.559066	0.528743
Heroin	0.038226	0.049237	0.055833	0.068662	0.094191	0.105396	0.121787	0.156625
Marijuana	0.506644	0.514837	0.490998	0.483837	0.471398	0.476172	0.478648	0.483462
Opioid	0.155213	0.177821	0.186006	0.220607	0.275012	0.303248	0.314244	0.339273
Cocaine	0.223397	0.194359	0.157602	0.146899	0.156809	0.131277	0.112446	0.106825
Other Drugs	0.000000	0.000000	0.119553	0.138593	0.154926	0.165516	0.177894	0.202654
Methamphetamine	0.087464	0.091235	0.096962	0.103556	0.112612	0.121995	0.141514	0.159936


```
In [27]: # Show a trend relative to other substances
draw_heatmap(df_grouped_year_T, 10, 7)
```



```
In [28]: # Draw a heatmap by County for Alcohol
draw_heatmap(
create_pivot_df(df_grouped,'Year','County Name', 'Alcohol'), 186, 20)
```




```
In [29]: #Opioid
draw_heatmap(create_pivot_df(df_grouped, 'Year', 'County Name', 'Opioid'), 186, 20)
```



```
In [30]: #Marijuana
draw_heatmap(create_pivot_df(df_grouped,'Year','County Name', 'Marijuana'), 186, 20)
```



```
In [31]: # Methamphetamine
draw_heatmap(create_pivot_df(df_grouped, 'Year', 'County Name', 'Methamphetamine'), 186, 20)
```




From the four heatmaps above, one could infer the following observations:

- while Alcohol remains the most commonly abused substance, its dominance is vaning as the prevalence of other substances, especially opioids and methamphetamine increases. At the same time, marijuana abuse remains fairly stable.

Suplimentary Analysis: Identify any years that differ from the majority

```
In [ ]: draw_heatmap_corr(create_pivot_df(df_grouped, 'Year', 'County Name', 'Alcohol'), 10,10)
```

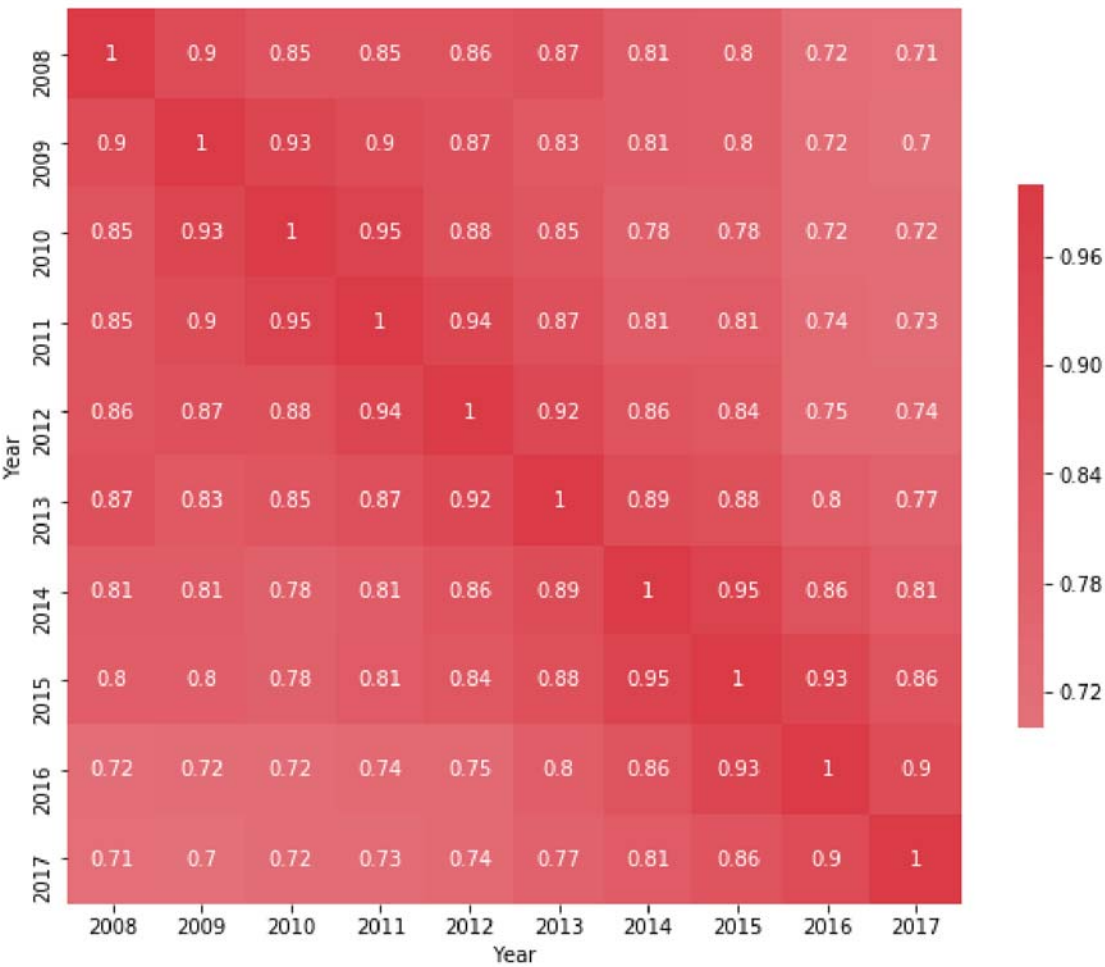
```
In [33]: #Marijuana
draw_heatmap_corr(create_pivot_df(df_grouped, 'Year', 'County Name', 'Marijuana'), 10,10)
```



```
In [34]: #Opioids
draw_heatmap_corr(create_pivot_df(df_grouped,'Year','County Name', 'Opioid'), 10,10)
```





```
In [35]: #Methamphetamine
draw_heatmap_corr(create_pivot_df(df_grouped, 'Year', 'County Name', 'Methamphetamine'), 1
0,10)
```




Beginning 2012 the usage patterns became very similar for the malority of the subjects that requested treatment in Indiana.

Suplimentary Analysis: Identify counties that differ from the majority



```
In [37]: draw_heatmap_corr(create_pivot_df(df_grouped, 'County Name', 'Year', 'Alcohol'), 93,93)
```





```
In [38]: draw_heatmap_corr(create_pivot_df(df_grouped, 'County Name', 'Year', 'Opioid'), 93,93)
```





```
In [39]: draw_heatmap_corr(create_pivot_df(df_grouped, 'County Name', 'Year', 'Marijuana'), 93,93)
```





```
In [40]: draw_heatmap_corr(create_pivot_df(df_grouped, 'County Name', 'Year', 'Other Drugs'), 93,93)
```




```
In [41]: draw_heatmap_corr(create_pivot_df(df_grouped, 'County Name', 'Year', 'Heroin'), 93,93)
```




```
In [42]: draw_heatmap_corr(create_pivot_df(df_grouped, 'County Name', 'Year', 'Methamphetamine'), 93, 93)
```



Suplimentary analisys: PCA

The principal component analysis (PCA) (<http://setosa.io/ev/principal-component-analysis/>) is the most basic dimensionality reduction method. To run the PCA we want to isolate only the numerical columns.

```
In [43]: from sklearn.decomposition import PCA  
pca = PCA()
```

Now I ran `fit()` method to identify principal components.

```
In [44]: pca_df_fitted = pca.fit(df_only_features)
print(pca.components_)
```

```
[[ -3.85284978e-01  1.74017452e-01 -1.44132829e-01  5.72548369e-01
   -2.69179141e-02 -9.45944737e-04  4.88667622e-01  2.68324235e-02
    4.82126552e-01]
 [ -4.18457857e-01 -5.05321504e-03  9.00800758e-01 -1.84296925e-02
   -5.53315489e-03  1.53651275e-03 -2.34787354e-02  1.09207683e-01
   -2.39850374e-02]
 [  7.11756112e-01 -2.43212736e-01  3.74058306e-01  9.21887106e-02
   -1.39977393e-01 -4.67858705e-02  3.29558300e-01 -2.17148697e-01
    3.29070826e-01]
 [  6.54654403e-02  2.33536949e-01  7.81258265e-02  1.44714164e-01
    6.66987977e-01 -5.86190546e-01 -6.12216977e-02 -3.43352728e-01
   -6.32252001e-02]
 [  7.12073209e-02  4.07014556e-01  9.88869180e-02  2.09251566e-01
    5.82755415e-02  6.54588384e-01 -1.35008442e-01 -5.53985972e-01
   -1.36726374e-01]
 [ -2.01654654e-01 -5.55680657e-01 -7.19214685e-02 -2.79439087e-01
    5.74989672e-01  3.75006655e-01  1.92924758e-01 -1.44138780e-01
    1.95080605e-01]
 [  3.46163005e-01  2.35182719e-01  8.23741063e-02  1.93313001e-01
    4.47977083e-01  2.91666275e-01 -1.35950867e-02  7.03277331e-01
   -1.29745391e-02]
 [  3.90868442e-04  5.68343537e-01  4.84514642e-03 -6.90996902e-01
    9.17264956e-03 -1.72477268e-03  2.33910579e-01  5.52881945e-03
    3.80333325e-01]
 [  1.39976189e-03  5.68623039e-02  4.14804014e-04 -7.54235243e-02
    1.18822447e-03  5.79425375e-04  7.33450549e-01  1.74171781e-03
   -6.73142890e-01]]
```

'-Other Drugs', Marijuana,Opioid+Alcohol, Marijuana+Cocaine, Marijuana+Alcohol, -Marijuana-Heroin, Marijuana+Methamphetamine

```
In [45]: print(pca.explained_variance_)
```

```
[0.52332706 0.25520259 0.20435398 0.13614426 0.12185475 0.1103172
 0.10171866 0.00989851 0.00189358]
```


```
In [46]: pca_df_fitted.explained_variance_ratio_
```

```
Out[46]: array([0.35729042, 0.17423414, 0.13951833, 0.0929496 , 0.08319374,
 0.07531672, 0.06944625, 0.006758 , 0.0012928 ])
```

```
In [47]: pca.components_.shape
```

```
Out[47]: (9, 9)
```

```
In [48]: plt.plot(np.cumsum(pca.explained_variance_ratio_))  
plt.xlabel('number of components')  
plt.ylabel('cumulative explained variance');
```



The first six components capture more than 95% of the variance in original dataset. This means that the PCA is not very effective on this dataset and six components will provide approximation for the rest of the dimensions.